

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



deeplearning.ai

Case Studies

Why look at
case studies?

Outline

Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception




deeplearning.ai

Case Studies

Classic networks

LeNet - 5



60K parameters.

$n_H, n_W \downarrow$ $n_C \uparrow$


conv pool conv pool fc fc output

Advanced: sigmoid/tanh ReLU

II, III.

↓

AlexNet




- Similar to LeNet, but much bigger.

- ReLU

- Multiple GPUs.

- Local Response Normalization (LRN)




~60M parameters




VGG - 16

CONV = 3x3 filter, s = 1, same

MAX-POOL = 2x2, s = 2




224x224x3



$n_h, n_w \downarrow$

$n_c \uparrow$

$\sim 138M$






deeplearning.ai


Case Studies

Residual Networks (ResNets)

Residual block



"short cut" / skip connection



$$\underline{z^{[l+1]}} = W^{[l+1]} \underline{a^{[l]}} + b^{[l+1]}$$


$$\underline{a^{[l+1]}} = g(\underline{z^{[l+1]}})$$

$$\underline{z^{[l+2]}} = W^{[l+2]} a^{[l+1]} + b^{[l+2]}$$


~~$$a^{[l+2]} = g(z^{[l+2]})$$~~

$$a^{[l+2]} = g(z^{[l+2]} + \underline{a^{[l]}})$$


Residual Network



Plain



ResNet






deeplearning.ai

Case Studies

Why ResNets work

Why do residual networks work?



ReLU. $\underline{a} \geq 0$

$$a^{[l+2]} = g(\underline{z}^{[l+2]} + \underline{a}^{[l]})$$

$$= g(\cancel{W^{[l+2]} a^{[l+1]} + b^{[l+2]}} + \underline{W_s a^{[l]}}) = g(a^{[l]}) = \underline{a^{[l]}}$$


$\uparrow$ 256 $\uparrow$ 128

$\mathbb{R}^{256 \times 128}$


If $W^{[l+2]} = 0, b^{[l+2]} = 0$

ResNet

Plain



ResNet





deeplearning.ai


Case Studies

Network in Network
and 1×1 convolutions

Why does a 1×1 convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6 \times 1$




$6 \times 6 \times 32$

*

2

=

32 $\rightarrow$ # filters.
 $n_c^{[l+1]}$




$1 \times 1 \times 32$

=

ReLU


Network in
Network

2	4	6	...		



$6 \times 6 \times \# \text{ filters}$

Using 1×1 convolutions






deeplearning.ai


Case Studies

Inception network motivation

Motivation for inception network





The problem of computational cost



Using 1×1 convolution

"bottleneck layer"



$$28 \times 28 \times 16 \times 192 = 2.4M$$

$$28 \times 28 \times 32 \times 5 \times 5 \times 16 = 10.0M$$

12.4M

120M




deeplearning.ai


Case Studies

Inception network

Inception module



Inception network







deeplearning.ai

Convolutional Neural Networks


MobileNet

Motivation for MobileNets

- Low computational cost at deployment
- Useful for mobile and embedded vision applications
- Key idea: Normal vs. depthwise-separable convolutions




Normal Convolution




$6 \times 6 \times 3$
 $n \times n \times n_c$

*



$3 \times 3 \times 3$
 $f \times f \times n_c$

=




$4 \times 4 \times 5$
 $n_{out} \times n_{out} \times n_c'$


$$\begin{aligned}
 \text{Computational cost} &= \frac{\text{\#filter params}}{3 \times 3 \times 3} \times \frac{\text{\# filter positions}}{4 \times 4} \times \text{\# of filters} \\
 &\rightarrow \underline{2160} = \uparrow \quad \quad \uparrow \quad \quad \uparrow
 \end{aligned}$$

Depthwise Separable Convolution


Normal Convolution



Depthwise Separable Convolution



Depthwise Convolution




Computational cost = #filter params x # filter positions x # of filters


$432 = \underbrace{3 \times 3}_{\text{filter size}} \times \underbrace{4 \times 4}_{\text{filter positions}} \times 3_{\text{filters}}$

Depthwise Separable Convolution


Depthwise Convolution



Pointwise Convolution



Pointwise Convolution




Computational cost = #filter params x # filter positions x # of filters


240 = $1 \times 1 \times 3 \times 4 \times 4 \times 5$

Depthwise Separable Convolution

Normal Convolution



Depthwise Separable Convolution



Cost Summary

Cost of normal convolution

2160

Cost of depthwise separable convolution

depthwise + pointwise
432 + 240 = 672

$$\frac{672}{2160} = 0.31 \leftarrow$$

$$= \frac{1}{\textcircled{n_c}} + \frac{1}{f^2}$$
$$\frac{1}{5} + \frac{1}{9}$$




$$= \frac{1}{512} + \frac{1}{3^2}$$

↑ ↑




~10 times cheaper

Depthwise Separable Convolution

Depthwise Convolution



Pointwise Convolution







deeplearning.ai

Convolutional Neural Networks

MobileNet Architecture



MobileNet

MobileNet v1



Pool, FC, Softmax


MobileNet v2



Pool, FC, Softmax


Bottleneck

MobileNet v2 Bottleneck




MobileNet


MobileNet v1



MobileNet v2



MobileNet v2 Full Architecture






deeplearning.ai

Convolutional Neural Networks

EfficientNet


EfficientNet

Baseline



A diagram showing a 3D volume represented by a blue oval. Three arrows point to the axes: 'r' for resolution, 'd' for depth, and 'w' for width.

Widener Rescuing








deeplearning.ai

Practical advice for using ConvNets

Transfer Learning

Transfer Learning





deeplearning.ai

Practical advice for using ConvNets

Data augmentation


Common augmentation method

Mirroring



y

Random Cropping




Rotation


Shearing

Local warping

...



Color shifting



Advanced:

PCA

ml-class.org


[AlexNet paper

["PCA color augmentation."

R B

G

Implementing distortions during training






deeplearning.ai

Practical advice for using ConvNets

The state of computer vision

Data vs. hand-engineering



Two sources of knowledge

- • Labeled data (x, y)
- • Hand engineered features/network architecture/other components



Tips for doing well on benchmarks/winning competitions

3-15 networks

→ $\hat{y}$

Ensembling

- Train several networks independently and average their outputs

Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1

+



4

+



1

+



4

Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset